

Dell | Cloudera Solution Deployment Guide v2.4

A Dell Deployment Guide for Apache™ Hadoop®
Crowbar v1.6

January 15, 2014



Table of Contents

Table of Contents.....	2
Figures.....	4
Tables.....	5
Trademarks.....	5
Notes, Cautions, and Warnings	6
Abbreviations.....	6
Overview.....	7
Summary	7
Network Setup.....	8
Managing Growth	8
Rack.....	8
Pod.....	8
Cluster	8
High-level Network Architecture	8
Network Overview.....	8
Layout.....	9
IP Addressing	10
Rack Awareness	10
Dell Cloudera Solution Hardware Configuration	12
Edge Node Configuration	12
Name Nodes Configuration.....	12
Data Nodes Configuration	13
Top of Rack Configuration	13
Dell Cloudera Solution Network Configuration	14
Dell Cloudera Solution Network Interconnects	15
Rack Configuration	16
Configuring the Dell™ Force10™ Network Solution	16
Single Rack Configuration	16
Stacking S60s.....	17
Uplinking the S60s	17
Server Gateway.....	18
Management Network	18
Multi-Rack Configuration	19
VRRP on S4810	20
VLT on S4810	21

Configuring the Force10 switch	22
Dell Cloudera Solution Deployment Process Overview	24
Dell Cloudera Solution Automated Software Installation	25
Admin Node Installation.....	25
Installing the Crowbar Software Framework	27
Troubleshooting Crowbar Admin Node Installations/Reinstallations	28
Duplicate Node IP Addresses	28
Virtual Media Impedes Installation	28
Editing the Network JSON.....	28
JSON Configuration by Section	30
How to Add a Public IP to a Node	31
Configuring the Network for Outbound Connectivity	31
Updating the public network definition	31
Allocating the IP Address	32
Verifying Master Node State	33
Data Node Installation	33
Troubleshooting a Hardware Installation Loop	33
Installing components	34
General Installation Process.....	34
Set CROWBAR Parameter	34
Obtain a proposal	34
Update a proposal	34
Committing a proposal	35
Modifying an active configuration	35
Installing Cloudera Manager	36
Hadoop Ecosystem Components.....	36
Dell Cloudera Solution Monitoring and Alerting.....	37
Appendix A: The Dell End User License Agreement	40
Appendix B: Network JSON Example	47
Appendix C: References	51
To Learn More	51

Figures

Figure 1: Dell Cloudera Solution Cabling 15

Figure 2: Single Rack View 16

Figure 3: Multi-Rack View 20

Figure 4: S4810 VLT interconnect..... 21

Figure 5: VLT Configuration on peer1 22

Figure 6: Dell | Cloudera Solution Deployment Process Overview 24

Figure 7: VMware Player Configuration for DVD..... 26

Figure 8: VMware Player Configuration for Network Adapter..... 27

Tables

Table 1: Dell Cloudera Solution Software Locations	7
Table 2: Dell Cloudera Solution Support Matrix	7
Table 3: Default Networks	8
Table 4: Master/Secondary (Admin) Name Nodes Network Connections	9
Table 5: Edge Nodes Network Connections	10
Table 6: Data Nodes Network Connections	10
Table 7: IP Addressing Schema	10
Table 8: Pod 1 IP Example Addressing Layout	10
Table 9: Pod 2 IP Example Addressing Layout	11
Table 10: Edge Node Hardware Configuration	12
Table 11: Master and Standby Name Node Hardware Configuration	12
Table 12: Data Node Hardware Configuration	13
Table 13: Network Top of Rack Configuration	13
Table 14: IP Scheme	14
Table 15: 60 Node Network	19
Table 16: Network Configuration Options	29
Table 17: Default Networks	29
Table 18: Network Parameters	29
Table 19: Range Map	30
Table 20: Crowbar Services	33
Table 21: Hadoop Ecosystem Components	36
Table 22: Components Monitored by Hadoop Monitoring Console	37

Trademarks

Reproduction of these materials is allowed under the Apache 2 license.

Information in this document is subject to change without notice.

© 2011-2014 Dell Inc. All rights reserved.

Dell, the DELL logo, and the DELL badge, PowerConnect, and PowerEdge are trademarks of Dell Inc. Cloudera, CDH, Cloudera Impala, and Cloudera Enterprise are trademarks of Cloudera and its affiliates in the US and other countries. Intel and Xeon are registered trademarks of Intel Corporation in the U.S. and other countries. Other trademarks and trade names may be used in this document to refer to either the entities claiming the marks and names or their products. Dell Inc. disclaims any proprietary interest in trademarks and trade names other than its own.

Other trademarks used in this text: Nagios®, Opscode Chef™, OpenStack™, Canonical Ubuntu™, and VMware™. Dell Precision™, OptiPlex™, Latitude™, PowerEdge™, PowerVault™, PowerConnect™, OpenManage™, EqualLogic™, KACE™, FlexAddress™ and Vostro™ are trademarks of Dell Inc. Intel®, Pentium®, Xeon®, Core™ and Celeron® are registered trademarks of Intel Corporation in the U.S. and other countries. AMD® is a registered trademark and AMD Opteron™, AMD Phenom™, and AMD Sempron™ are trademarks of Advanced Micro Devices, Inc. Microsoft®, Windows®, Windows Server®, MS-DOS® and Windows Vista® are either trademarks or registered trademarks of Microsoft Corporation in the United States and/or other countries. Red Hat Enterprise Linux® and Enterprise Linux® are registered trademarks of Red Hat, Inc. in the United States and/or other countries. Novell® is a registered trademark and SUSE™ is a trademark of Novell Inc. in the United States and other countries. Oracle® is a registered trademark of Oracle Corporation and/or its affiliates. Citrix®, Xen®, XenServer® and XenMotion® are either registered trademarks or trademarks of Citrix Systems, Inc. in the United States and/or other countries. VMware®, Virtual SMP®, vMotion®, vCenter®, and vSphere® are registered trademarks or trademarks of VMware, Inc. in the United States or other countries.

January 15, 2014

Notes, Cautions, and Warnings



A **NOTE** indicates important information that helps you make better use of your computer.



A **CAUTION** indicates potential damage to hardware or loss of data if instructions are not followed.



A **WARNING** indicates a potential for property damage, personal injury, or death.

Abbreviations

Abbreviation	Definition
BMC	Baseboard Management Controller
CDH	Cloudera Distribution for Hadoop
DBMS	Database management system
EDW	Enterprise data warehouse
EoR	End-of-row switch/router
HDFS	Hadoop Distributed File System
IPMI	Intelligent Platform Management Interface
NIC	Network interface card
LOM	Local area network on motherboard
OS	Operating system
ToR	Top-of-rack switch/router

Overview

Summary

This deployment guide for the Dell™ | Cloudera™ Solution describes the steps to install the solution on a predefined hardware and network configuration as specified in the current ***Dell | Cloudera Solution Reference Architecture*** document. It covers the steps required to prepare hardware platforms for the deployment of Cloudera Manager or Cloudera Hadoop (CDH). For deployment of Cloudera Manager, use the ***Dell | Cloudera Solution User's Guide***.

Table 1: Dell | Cloudera Solution Software Locations

Daemon	Primary Location	Secondary Location
JobTracker	Master Name Node	
TaskTracker	Data Node(x)	
NameNode	Master Name Node	Secondary Name Node
Operating System Provisioning	Admin Node	
Chef	Admin Node	
Yum Repositories	Admin Node	
Cloudera Management Suite	Edge Node(x)	
Zookeeper	Data Node(x)	
HMaster	Master Name Node	
RegionServer	Data Node(x)	
Crowbar Admin	Admin Node	

Table 2: Dell | Cloudera Solution Support Matrix

Category	Component	Version	Available Support
Operating System	Red Hat Enterprise Linux	6.4	Red Hat Linux support
Operating System	CentOS	6.4	Dell Hardware support
Java Virtual Machine	Sun Oracle JVM	Java 6	N/A
Hadoop	Cloudera Distribution for Apache Hadoop (CDH)	4.5	Cloudera support
Hadoop	Cloudera Manager	4.8	Cloudera support
Hadoop	Impala	1.2	Cloudera support
Hadoop	Cloudera Search	1.1	Cloudera support

Network Setup

The network configuration assumes a flat L2 wiring. All network connections should be accessible at that layer. Where isolation between different logical networks is required, VLANs are used.

Managing Growth

The system architecture is organized into three components, for sizing as the Hadoop environment grows. From smallest to largest, they are:

- Rack
- Pod
- Cluster

Each has specific characteristics and sizing considerations. You can scale the environment by adding additional capacity as needed, without the need to replace any existing components.

Rack

A rack is the smallest component in a Hadoop environment, and consists of all of the power, network cabling, and two Ethernet switches required to support up to 20 data nodes. These nodes should utilize their own power connectivity and data center space – separate from other racks – and be treated as a fault zone.

Pod

A pod is a single set of stacked Ethernet switches. For the Dell | Cloudera Reference Architecture, both the maximum and minimum are six. A pod consists of the administration and operation infrastructure to support three racks.

Cluster

A cluster is a set of greater than one pod, up to a maximum of 12 pods. A cluster is a set of Hadoop nodes that share the same Network Node and management tools for operating the Hadoop environment.



Please see the current *Dell | Cloudera Solution - Reference Architecture Guide* for more detailed information

High-level Network Architecture

Network Overview

The Dell | Cloudera solution implements at a minimum four distinct, separate VLANs:

Table 3: Default Networks

Usage	Description	Default reserved VLAN tag	Tagged
Management vLAN	Used for administrative functions such as Crowbar node installation, TFTP booting, DHCP assignments, system logs, backups, and other monitoring. There is only one vLAN set up for this function and it is spanned across the entire network.	100	Not tagged
BMC vLAN	Used for connecting to the BMC of each node.	100	Not tagged
Production vLAN	Used by the Hadoop system to handle traffic between all nodes for HDFS operations, MapReduce jobs, and other Hadoop traffic.	200	802.1q Tagged

Usage	Description	Default reserved vLAN tag	Tagged
Edge vLANs	Used for connections to devices external to the Hadoop cluster; these include externally visible services such as load balancers and web servers. Use one or many of these networks, dependent on the need to segregate traffic among groups of servers	300	802.1q Tagged

 The management and BMC networks are expected to be in the same L2 network.

All servers in a Hadoop cluster are tied together using TCP/IP networks. These networks form a data interconnect across which individual servers pass data back and forth, return query results, and load/unload data. These networks are also used for management.

The Admin Node manages all the cluster nodes. It assigns the other nodes IP addresses, PXE boots them, configures them, and provides them the necessary software for their roles. To provide these services, the Admin Node runs the **Dell™** Crowbar software framework, Chef, DHCP, TFTP, NTP, and other services, and this must be the only DHCP server visible to the infrastructure and data nodes. Details follow:

- **DHCP server—assigns and manages IPs for the compute and storage nodes**
- **NTP server (Network Time Protocol server)—makes sure all nodes are keeping the same clock**
- **TFTP server—PXE boots compute and storage nodes with a Linux kernel; the TFTP server services any PXE boot request it receives with its default options.**
- **DNS server—manages the name resolution for the nodes and can be configured to provide external name forwarding.**

Due to the nature of the different software used, the network is set up as flat as possible using a dedicated BMC port and bonded LOMs. If the Crowbar software framework is used to deploy the cluster, it manages all networks, and comes out of the box preconfigured to allow the initial configuration to come up quickly by predefining the admin, public, and BMC networks.

The Crowbar network configuration can be customized to better map to site-specific networking needs and conventions. These changes include adding additional vLANs, changing vLAN mappings, and teaming NICs.

Layout

Due to the nature of Crowbar's network layout, addresses are assigned to a whole network based upon interface, Network Type (Production, Management, and External) and teaming type.

Table 4: Master/Secondary (Admin) Name Nodes Network Connections

Interface	Network Type	Teaming Type
BMC	Management LAN	Single
LOM1	Production LAN	Teamed
LOM2	Production LAN	Teamed
Eth1	Production LAN	Teamed
Eth2	Management LAN	Single

Table 5: Edge Nodes Network Connections

Interface	Network Type	Teaming Type
BMC	Management LAN	Single
LOM1	Production LAN	Teamed 1
LOM2	Production LAN	Teamed 1
Eth1	Edge LAN	Teamed 2
Eth2	Edge LAN	Teamed 2

Table 6: Data Nodes Network Connections

Interface	Network Type	Teaming Type
BMC	Management LAN	Single
LOM1	Production LAN	Teamed 1
LOM2	Production LAN	Teamed 1

IP Addressing

The IP address can be assigned in this fashion, using large subnets to support many machines on the production network. The management network is a Class C network with 254 IP addresses. The Production network is what is known as a /23 with 512 IP addresses. In each network, the first 10 IP addresses are reserved for switches, routers, and firewalls.


 Each network's ".1" address is reserved for the network gateway.

Table 7: IP Addressing Schema

LAN	Network	Subnet	Gateway	Reserved
Management LAN	172.16.0.0	255.255.255.0	172.16.0.1	0.1 – 0.10
Production LAN	172.16.2.0	255.255.254.0	172.16.2.1	2.1-2.20
Name Nodes	DHCP Allocated			
Data Nodes	DHCP Allocated			
External LAN	TBD by Customer			

Rack Awareness

With the network set up using Top of Rack (ToR) switches, Rack Awareness can be programmed using the Chef information about which switch the LOM1 is plugged into. A simple script has been added to the Hadoop configuration to pull the information out of Chef, and then use it for Rack Awareness.

Table 8: Pod 1 IP Example Addressing Layout

Network: 172.16.0.0	Netmask: 255.255.252.0
Multicast: 172.16.0.0	Broadcast 172.16.3.255

Pod	Rack Number	Network	Server Type	IP Range	Subnet Mask	Gateway
1	1	Production	Data Node	172.16.0.1-42	255.255.252.0	172.16.0.1
1	2	Production	Data Node	172.16.1. 1-42	255.255.252.0	172.16.0.1
1	3	Production	Data Node	172.16.2. 1-42	255.255.252.0	172.16.0.1
1		Production	Master Name Node	172.16.3.1-19	255.255.252.0	172.16.0.1

Pod	Rack Number	Network	Server Type	IP Range	Subnet Mask	Gateway
1		Production	Secondary Name Node	172.16.3.20-30	255.255.252.0	172.16.0.1
1		Production	Edge Node	172.16.3.41-50	255.255.252.0	172.16.0.1
1	1	BMC	Data Node	172.16.0.200-242	255.255.252.0	172.16.0.1
1	2	BMC	Data Node	172.16.1.200-242	255.255.252.0	172.16.0.1
1	3	BMC	Data Node	172.16.2.200.242	255.255.252.0	172.16.0.1
1		BMC	Master Name Node	172.16.3.201-219	255.255.252.0	172.16.0.1
1		BMC	Secondary Name Node	172.16.3.220-230	255.255.252.0	172.16.0.1
1		BMC	Edge Node	172.16.3.231-250	255.255.252.0	172.16.0.1

Table 9: Pod 2 IP Example Addressing Layout

Network: 172.16.0.0 Netmask: 255.255.252.0

Multicast: 172.16.0.0 Broadcast: 172.16.3.255

Pod	Rack Number	Network	Server Type	IP Range	Subnet Mask	Gateway
2	1	Production	Data Node	172.16.4.1-42	255.255.252.0	172.16.4.1
2	2	Production	Data Node	172.16.5. 1-42	255.255.252.0	172.16.4.1
2	3	Production	Data Node	172.16.6. 1-42	255.255.252.0	172.16.4.1
2		Production	Master Name Node	172.16.7.1-19	255.255.252.0	172.16.4.1
2		Production	Secondary Name Node	172.16.7.20-30	255.255.252.0	172.16.4.1
2		Production	Edge Node	172.16.7.41-50	255.255.252.0	172.16.4.1
2	1	BMC	Data Node	172.16.4.200-242	255.255.252.0	172.16.4.1
2	2	BMC	Data Node	172.16.5.200-242	255.255.252.0	172.16.4.1
2	3	BMC	Data Node	172.16.6.200.242	255.255.252.0	172.16.4.1
2		BMC	Master Name Node	172.16.7.201-219	255.255.252.0	172.16.4.1
2		BMC	Secondary Name Node	172.16.7.220-230	255.255.252.0	172.16.4.1
2		BMC	Edge Node	172.16.7.231-250	255.255.252.0	172.16.4.1
2		External	Edge Node	TBD by Customer	TBD	TBD

Dell | Cloudera Solution Hardware Configuration

Edge Node Configuration

Table 10: Edge Node Hardware Configuration

Component	Setting	Parameter
BIOS	Boot Order	LOM 1 PXE
		Internal Boot Device PERC H710 LUN 0
	PXE Boot LOM 1	Enable
	PXE Boot LOM 2	Disable
	C-State	Disable
PERC H710 BIOS	RAID	Enabled
	LUN 0	Disk 0-5 RAID 10
	Boot Order	LUN 0

Name Nodes Configuration

Table 11: Master and Standby Name Node Hardware Configuration

Component	Setting	Parameter
BIOS	Boot Order	LOM 1 PXE
		Internal Boot Device PERC H710 LUN 0
	PXE Boot LOM 1	Enable
	PXE Boot LOM 2	Disable
	C-State	Disable
PERC H710 BIOS	RAID	Enabled
	LUN 0	Disk 0-5 RAID 10
	Boot Order	LUN 0

Data Nodes Configuration

Table 12: Data Node Hardware Configuration

Component	Setting	Parameter
BIOS		
	Boot Order	LOM 1 PXE Internal Boot Device
	PXE Boot LOM 1	Enable
	PXE Boot LOM 2	Disable
LSI 9207 Controller BIOS		
	Number of INT13 devices	24
	IO Timeout for Block Devices	10
	IO Timeout for Block Devices (Removable)	10
	IO Timeout for Sequential Devices	10
	IO Timeout for Other Devices	10
	LUN's to Scan for Block Devices	All
	LUN's to Scan for Block Devices (Removable)	All
	LUN's to Scan for Sequential Devices	All
	LUN's to Scan for Other Devices	All
	Removable Media Support	None
PERC H710 Controller BIOS		
	RAID	Enabled
	LUN0	Disk0 RAID0
	LUN1	Disk1 RAID0

	LUN23	Disk23 RAID0
	Boot Order	Disk 0 Disk 1

- Data nodes can be configured with either the LSI 9207 or the PERC H710 disk controller. The LSI 9207 is recommended for new deployments. The PERC H170 is supported as an alternative, primarily for compatibility with existing clusters.

Top of Rack Configuration

Table 13: Network Top of Rack Configuration

Setting	Parameter	Ports
Spanning-Tree	Enable	ALL
Port-Fast	Enable	ALL
Flow-Control	Enable	ALL

Dell | Cloudera Solution Network Configuration


Table 14: IP Scheme

A	B	C	D	Use
First POD				
172	16	0/22		Rack Number
			1-42	Data Node[XX] bond0, by Rack Unit
		4/22		Rack Number (1xx)
			200-242	Data Node [XX] BMC, by Rack Unit
172	16	3	1-19	Master Node[XX]
		3	20-30	Data Node[XX]
		3	41-50	Edge Node[XX]
172	16	7	1-19	Master Node[XX]
		7	20-30	Master Node[XX]
		7	41-50	Edge Node[XX]

- All Master Nodes will be addressed in the first Pod only. Additional Pods will not contain additional Master Nodes.
- Master Nodes running Zookeeper-related services will be distributed among Pods for larger deployments. Please consult with your Dell sales team when designing your solution.

Dell | Cloudera Solution Network Interconnects

Figure 1: Dell Cloudera Solution Cabling



- Notes On the Infrastructure Nodes:
- 10GbE ports 1 and 3 are bonded
 - 10GbE ports 2 and 4 are bonded.
 - all ports are not used.

Rack Configuration

Configuring the Dell™ Force10™ Network Solution


Single Rack Configuration

With 1G nodes, Dell Force10 recommends using Dell™ Force10™ S60 ToR switches in the rack. Each rack could have a maximum of 20 servers. Each rack has two ToR S60 switches that are stacked, and this stack connects to **both of the Dell™ Force10™ S4810 switches**. The S60 stack offers a single switch view to the servers. Each Data Node has two data 1GbE NIC ports. It forms a LAG of two ports with one port on each switch in the stack, thereby offering high availability and redundancy.

For the 10G servers, the pair of ToR switches are Dell Force10 S4810, that run high availability feature called the Virtual Link Trunking (VLT). This feature allows the servers to terminate their LAG interfaces into two different switches instead of one. This allows HA as well as active-active bandwidth utilization. This feature gives redundancy within the rack if one switch fails or needs maintenance. The uplink to the aggregation pair is a LAG, from each ToR, of 80G bandwidth. This is achieved using 2 x 40G interfaces in a LAG connecting to the aggregation pair. Therefore from each rack there is a collective bandwidth of 160G available.

In addition, each 10G node will have a single the On-Board NIC, eth0, connected to either a Dell Force10 S55, S60, or Dell PowerConnect 6248, along with the BMC of each node. The On-Board NIC will be used for PXE booting and initial configuration of each node.

Figure 2: Single Rack View



Use the switch configuration guide (manual) for the initial configurations. Examples of these are: enabling the interfaces ('no shut'), configuration of IPs on the management interfaces, enabling SSH (telnet is enabled by default), and authorization details.

Stacking S60s

The following configuration helps stack the two S60 switches together within the rack. This configuration assumes the stacking module in both S60 switches is in module 0 (IO-facing side) and the 10G uplink module is in slot 1 (power supply and fan side).

Connect port 49 on module 0 (IO side from the left) to port 49 of the second S60 and similarly connect port 50 on both switches. The stack is automatically detected and formed without a user configuration. Using the CLI command 'show system brief' verify that the stacking module is detected by the S60.

When you are adding units to a stack, you can either:

- **Allow FTOS to automatically assign the new unit a position in the stack, or**
- **Manually determine each unit's position in the stack by configuring each unit to correspond with the stack before connecting it.**

Three configurable system variables affect how a new unit joins a stack: priority, stack number, and provision.

After the new unit loads, it synchronizes its running and startup configurations with the stack.

```
TOR-Rack1#stack-unit renumber
TOR-Rack1(config)# stack-unit priority <higher priority determines primary role>
```

After connecting the switches together run the following command to check the status of the stack.

```
TOR-Rack1#show system brief

Stack MAC : 00:01:e8:d5:ef:81

-- Stack Info --

Unit      UnitType   Status   ReqTyp   CurTyp   Version   Ports
-----
0         Standby    online   S60      S60      8.3.3.7   52
1         Management online   S60      S60      8.3.3.7   52
```

Uplinking the S60s

The following configuration helps create configurations for the uplink of the stack. This configuration assumes the 10G uplink module is in slot 1 (power supply and fan side). The uplink ports are going to be numbered 0/51,0/52 and 1/51,1/52 respectively. All four 10G interfaces would be part of a single LAG or port-channel. The following illustrates that.

```
# Put the user ports in the switchport mode
TOR-Rack1(config)# interface range gigabitethernet 0/1 - 47
TOR-Rack1(config-if-range-gi-0/1-47)# no shutdown
TOR-Rack1(config-if-range-gi-0/1-47)#switchport
TOR-Rack1(config-if-range-gi-0/1-47)#end

# Repeat the same for ports on the second unit
TOR-Rack1(config)# interface range gigabitethernet 1/1 - 47
<snip>...
```

```

# Create port-channel of the 4 10G ports. The example below shows it for 1
port.
# Repeat the same configs for other 10G ports 0/52,1/51 and 1/52.

TOR-Rack1(conf)#interface GigabitEthernet 0/51
TOR-Rack1(conf-if-gi-3/15)#no shutdown
TOR-Rack1(conf-if-gi-3/15)#port-channel-protocol lacp
TOR-Rack1(conf-if-gi-3/15-lacp)#port-channel 1 mode active

# Change the defaults on the port-channel that gets created automatically
# From the above commands.

TOR-Rack1(conf)#interface port-channel 1
TOR-Rack1(conf-if-po-1)#no shutdown
TOR-Rack1(conf-if-po-1)#switchport

# Add the Data ports 0 through 30 and the port-channel 1 to vlan 100

TOR-Rack1#config
TOR-Rack1 (conf)#int vlan 100
TOR-Rack1 (conf-if-vlan)#tagged po 1
TOR-Rack1 (conf-if-vlan)#untagged gi 0/0-21
TOR-Rack1 (conf-if-vlan)#untagged gi 1/0-21
TOR-Rack1 (conf-if-vlan)#show conf
!
interface Vlan 100
no ip address
tagged Port-channel 1
untagged gi 0/0-21
untagged gi 1/0-21

TOR-Rack1#config
TOR-Rack1 (conf)#int vlan 300
TOR-Rack1 (conf-if-vlan)#tagged po 1
TOR-Rack1 (conf-if-vlan)#untagged gi 0/29-41
TOR-Rack1 (conf-if-vlan)#show conf
!
interface Vlan 300
no ip address
tagged Port-channel 1
untagged gi 0/29-41

```

So far the configuration is sufficient to link the nodes to the ToR switches, stacking the ToR and uplinks from ToR.

The uplink port-channel links are all active and forward traffic to the aggregate switches. Each **flow**, a unique combination of a source and a destination, gets hashed internally and gets load-balanced across the port-channel.

Server Gateway

The nodes in a rack have a single virtual IP as their gateway for routing purposes. The VRRP protocol runs on the aggregation S4810 switches. It does not need any configuration on the ToR. The VRRP master owns the virtual IP and does the routing but the combination of VLT and VRRP makes it certain that backup also routes or switches the traffic if it has a path in its forwarding table. This is an active-active brained capability where routing is independent of which switch owns the virtual IP.

Management Network

The BMC ports from all the nodes connect to the same ToR switches as the data ports. However the management vLAN is separate from the data vLAN. Ports 0 to 30 on the ToR are reserved for data connections and 31 to 48 are configured for the management network. This is achieved by creating a separate vLAN on the ToR and adding all the management ports as part of that vLAN.

```

TOR-Rack1(conf)#int vlan 300
TOR-Rack1(conf-if-vlan)#tagged po 1
TOR-Rack1(conf-if-vlan)#untagged gi 0/31-47
TOR-Rack1(conf-if-vlan)#untagged gi 1/31-47

```

Multi-Rack Configuration


Once the single rack is deployed from the server and network perspective we can take a look at the multi-rack view and then move on to configure the aggregation switches that connect the racks together. This section shows the S4810 switch aggregating the clusters together to enable inter-rack traffic as well as the management network. As we saw there are two separate vLANs for data and management; all port-channels on the S4810 and ToR are tagged in these two vLANs.

The following table shows the network inventory details in a full cluster of three racks.

Table 15: 60 Node Network

Total Racks	3 (15-20 nodes per rack)
Top of Rack Switch	6 S60 (2 per rack)
Pod-interconnect Switch	2 S4810
Server	2RU R720/R720xd
Over-subscription at ToR	1:1
Modules in Each ToR	1x 12-2port Stacking, 1x 10G -2 port uplink

Figure 3: Multi-Rack View



VRRP on S4810

The following configuration shows a sample VRRP configuration on the S4810 switches. This configuration is created on the vLAN interfaces of the S4810. Since there is only a single vLAN 100 in the cluster of three racks, a single instance of this configuration is needed.

```

Force10_VLTpeer1(conf)#int vlan 100
Force10_VLTpeer1(conf-if-gi-1/1)#vrrp-group 100
Force10_VLTpeer1(conf-if-gi-1/1-vrid-111)#virtual-address 10.10.10.1
#One or more these virtual IP addresses can be configured, which can be used
#as the unique gateway per rack or cluster.
Force10_VLTpeer1(conf-if-gi-1/1-vrid-111)# priority 125
# Priority from 1-255 can be used to determine which switch owns the VIP and
#becomes the VRRP master.


# Repeat the same configuration on the second VLT peer, except for a different
priority.

```

VLT on S4810

The second part of configuration involves the pod-interconnect switches that run VLT with each other.

Figure 4: S4810 VLT interconnect



With the following steps we will configure VLT on the pair of S4810 switches that interconnect the racks. To configure virtual link trunking, you must create a VLT domain, configure a backup link and interconnect trunk, and connect the peer switches in a VLT domain to an attached access device (switch or server). But first RSTP should be configured, as a best practice, on the S4810 as well as the S60 switches.

```

Force10_VLTpeer1(conf)#protocol spanning-tree rstp
Force10_VLTpeer1(conf-rstp)#no disable
Force10_VLTpeer1(conf-rstp)#bridge-priority 4096

#Repeat the same on VLTPeer2 with a different bridge priority to make it the
root.

Force10_VLTpeer2(conf-rstp)#bridge-priority 0

```

The next figure shows a sample configuration on VLT. The VLT works over a primary link and a backup link. Therefore this configuration consists of configuring the IP connectivity details of each switch. In addition each port-channel to the layer-2 switch, S60 stack in this case, gets a configuration specifying the port-channel that acts as the ICL link. In absence of a direct path to the destination, the ICL link would carry the traffic to the peer. The backup link is only for heartbeat status of the peer; no data traffic flows over it.

Figure 5: VLT Configuration on peer1

```

Forcel0_VLTpeer1(conf)#vlt domain 999
Forcel0_VLTpeer1(conf-vlt-domain)#peer-link port-channel 100
Forcel0_VLTpeer1(conf-vlt-domain)#back-up destination 10.11.206.35
Forcel0_VLTpeer1(conf-vlt-domain)#exit

Forcel0_VLTpeer1(conf)#interface ManagementEthernet 0/0
Forcel0_VLTpeer1(conf-if-ma-0/0)#ip address 10.11.206.23/16
Forcel0_VLTpeer1(conf-if-ma-0/0)#no shutdown
Forcel0_VLTpeer1(conf-if-ma-0/0)#exit

Forcel0_VLTpeer1(conf)#interface port-channel 100
Forcel0_VLTpeer1(conf-if-po-100)#no ip address
Forcel0_VLTpeer1(conf-if-po-100)#channel-member fortyGigE 0/56,60
Forcel0_VLTpeer1(conf-if-po-100)#no shutdown
Forcel0_VLTpeer1(conf-if-po-100)#exit

Forcel0_VLTpeer1(conf)#interface port-channel 110
Forcel0_VLTpeer1(conf-if-po-110)#no ip address
Forcel0_VLTpeer1(conf-if-po-110)#switchport
Forcel0_VLTpeer1(conf-if-po-110)#channel-member fortyGigE 0/52
Forcel0_VLTpeer1(conf-if-po-110)#no shutdown
Forcel0_VLTpeer1(conf-if-po-110)#vlt-peer-lag port-channel 110
Forcel0_VLTpeer1(conf-if-po-110)#end

Forcel0_VLTpeer1# show vlan id 10
Codes: * - Default VLAN, G - GVRP VLANs, P - Primary, C - Community, I - Isolated
Q: U - Untagged, T - Tagged
  x - Dot1x untagged, X - Dot1x tagged
  G - GVRP tagged, M - Vlan-stack, H - Hyperpull tagged

  NUM      Status      Description                               Q Ports
  10       Active
                                     U Po110 (Fo 0/52)
                                     T Po100 (Fo 0/56,60)

```

Enable VLT and create a VLT domain with a backup-link and interconnect trunk

Configure the backup link

Configure the VLT trunk interconnect

Configure the port channel to an attached device

Verify that the port channels used in the VLT domain are assigned to the same VLAN

Configuring the Force10 switch

1. Use a serial communication hyperterminal (e.g. Minicom) to configure the switch. The following are the instructions for using Minicom:

```
# minicom -s
- Serial port settings: /dev/ttyUSB0, 9600,n,8,1
- Modem and Dialing:
  o Delete init, dial and hangup lines
- Save settings as df1
```

```
# minicom
```

```
force10> enable
```

```
# configure
```

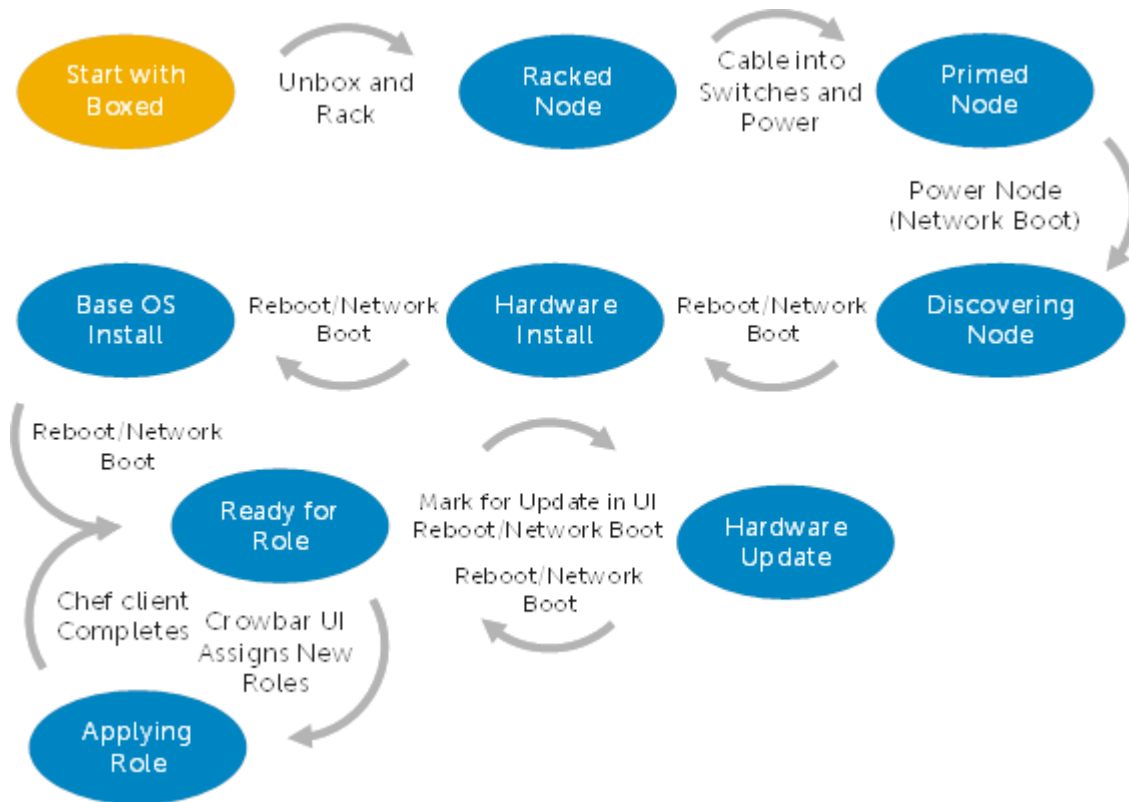
```
interface range gigabitethernet 0/0 - 46
```

```
no ip address
```

```
switchport
spanning-tree rstp edge-port
no shutdown
exit
interface range gigabitethernet 1/0 – 46
no ip address
switchport
spanning-tree rstp edge-port
no shutdown
exit
Int vlan 100
description Production
no ip address
untagged gigabitethernet 0/0-30
untagged gigabitethernet 1/0-30
no shutdown
exit
int vlan 300
description BMC
no ip address
untagged gigabitethernet 0/31-39
no shutdown
exit
```

Dell | Cloudera Solution Deployment Process Overview


Figure 6: Dell | Cloudera Solution Deployment Process Overview



Dell | Cloudera Solution Automated Software Installation

Admin Node Installation

To use the Crowbar software framework, you must first install an Admin Node. Installing the Admin Node requires installing the base operating system, optionally customizing the Crowbar configuration, and installing Crowbar itself.

 The Crowbar admin node acts as a DHCP server. To avoid problems with installation and node discovery, you should ensure that there are no other Crowbar admin nodes or DHCP servers on the network.

In addition, Crowbar supports only one (1) Top-Level Domain (TLD).

The following is required to bootstrap the Admin Node by PXE booting:

1. The user is required to make the physical arrangements to connect a VM to the network in such a way that the Admin Node (when there is one) can PXE boot from it. A network crossover cable may be required.
2. All BIOS and RAID configuration for the Crowbar Admin Node will need to be completed manually, prior to the installation from the Crowbar ISO image.
3. A VM image provides an initial TFTP/DHCP/Boot server. A VMware Player (free download from VMware) is required to execute it.

In preparation for running VMware Player on a particular machine, please make sure that:

- **Support for Intel VT is enabled in BIOS**
- **There is only one NIC enabled (turn off the wireless NIC if there is one and leave only the wired NIC enabled)**

Procedure:

1. Install VMware Player on a laptop.
2. Open the VMware machine configuration distributed with the Crowbar software framework. (e.g., Crowbar_Installer-1.3.tgz)
3. Edit the machine settings (see figures that follow) and ensure that:
 - **The CD/DVD drive is mounting the Crowbar ISO distribution**
 - **The Network adapter is configured to use Bridged Networking**
4. Obtain the ISO of Crowbar (from your Dell Account Representative) and configure VMware Player to mount it as a DVD in the VM.
5. Plug the crossover cable into eth0 of the server and your network port on the laptop.
6. Start the VMware Player and configure it to use the network port.
7. Power on the Admin Node, and ensure that:
 - **It is set up to boot from the hard disk for subsequent boots**
 - **The first boot is a network boot**

The machine will obtain its image from the VMware Player VM and start the installation process.

 PXE booting is not enabled by default on all NICs. If you change NICs you may need to manually configure them, to enable PXE booting.

Figure 7: VMware Player Configuration for DVD

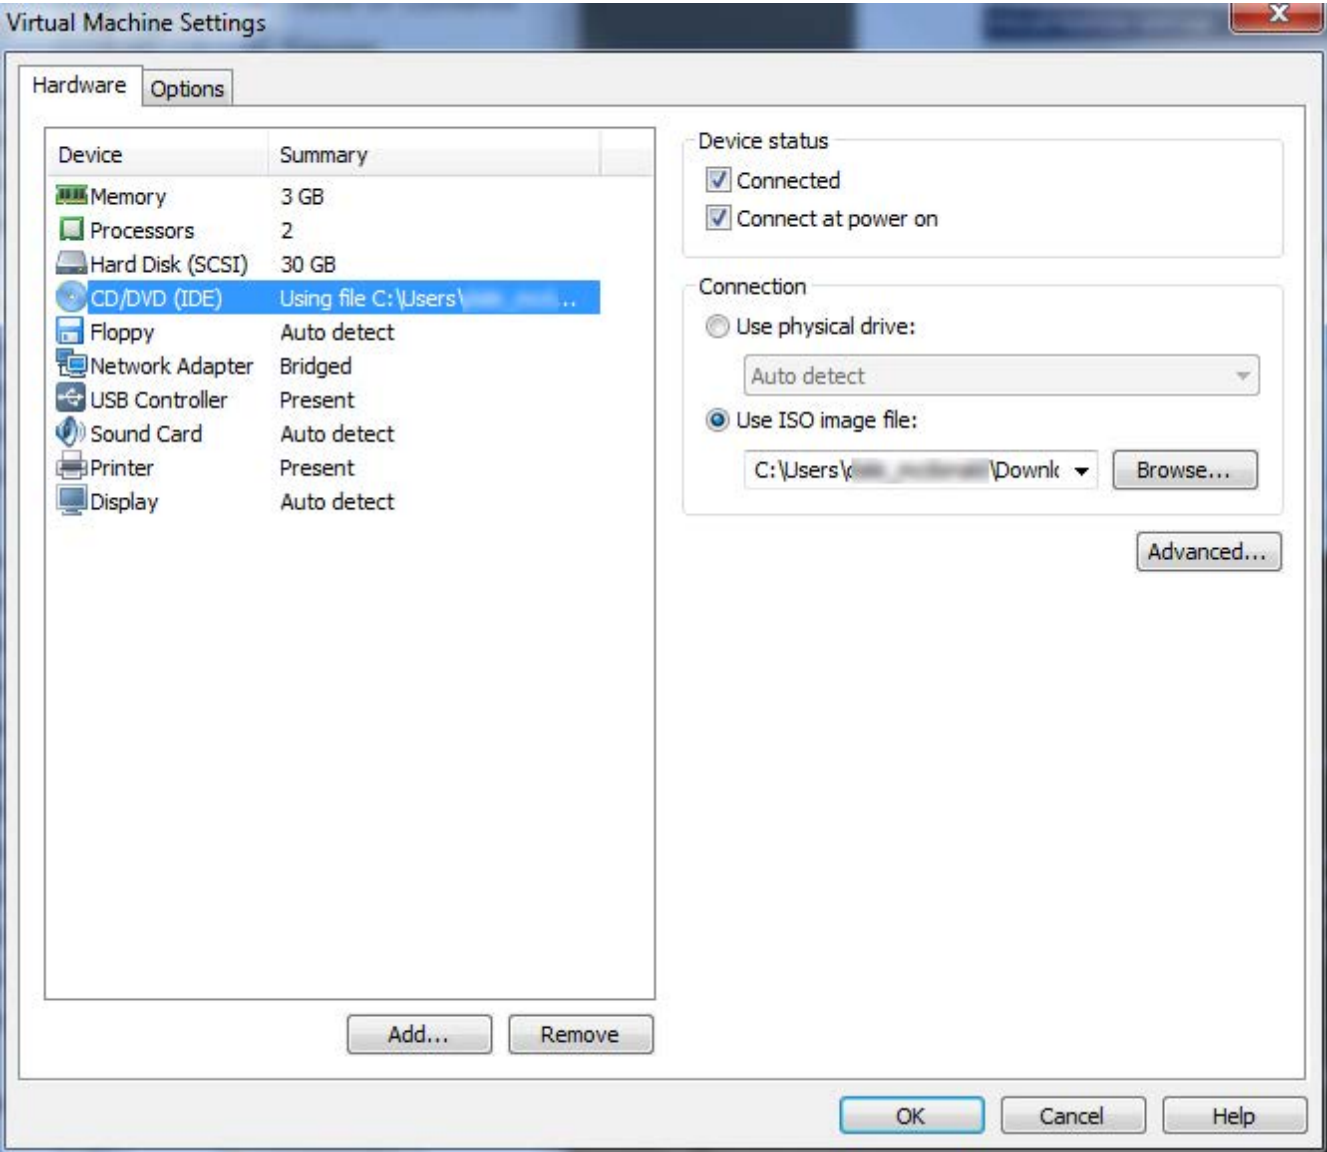



Figure 8: VMware Player Configuration for Network Adapter



Installing the Crowbar Software Framework

The image installed in the previous steps includes all the required Crowbar components. Before actually installing the Crowbar software framework, there is the opportunity to customize the installation to fit into the deployment environment. The steps below assume default configuration.

To install the Crowbar software framework:

1. Log onto the Admin Node. The default username is `root`, password: `crowbar`.
2. If necessary, edit the file `/opt/dell/barclamps/network/chef/data_bags/crowbar/bc-template-network.json` to customize the network information for the deployment. A detailed description of how to edit the network JSON can be found in the next section.
 - a. The networks cannot be reconfigured once the system is installed.
3. `cd /tftpboot/redhat_dvd/extra`
4. `./install admin.your.cluster.fqdn`
 - a. where `admin.your.cluster.fqdn` is the fully qualified domain name of the admin machine, for example `admin.dell.com`

- b. The installation is run in a `screen` process. To monitor progress during the installation, attach to the screen using the `screen -r` command.
5. Verify in `/var/log/install.log` that the install completed successfully. Look for "*Script done on <date_time>*"
6. Verify that no installation errors occurred.
7. Reboot the Crowbar admin node.



Because there are many dependencies some transient errors might be visible on the console. This is expected.

Troubleshooting Crowbar Admin Node Installations/Reinstallations

This section provides guidance to help you troubleshoot Crowbar admin node installation or reinstallation issues.

Duplicate Node IP Addresses

- **Problem:** If previously-installed nodes with IP addresses in the BMC network range are turned off, and new nodes are subsequently allocated, the new nodes may be configured with the IP addresses of the powered-off nodes.
- **Workaround:** To avoid this, ensure that the BMC IP addresses on all nodes are unique.

Virtual Media Impedes Installation

- **Problem:** iDRAC virtual media is listed as attached, preventing installation/reinstallation.
- **Workaround:** Ensure that no virtual media, including USB thumb drives, are attached. Then, proceed with the installation/reinstallation.

Editing the Network JSON

The JSON is located at `/opt/dell/barclamps/network/chef/data_bags/crowbar/bc-template-network.json`. The file should be edited before the `install` command is run to create your Admin Node.

The information you will need:

1. VLAN ID for each of the VLANs used by Crowbar
2. Network subnet for each VLAN
3. Ranges for each network (i.e., host, dhcp, switch, admin)
4. Netmask for each VLAN
5. Gateway for the public network and possibly the BMC ranges

See Appendix B: Network JSON Example.



The biggest error for many users is putting a comma at the end of the last statement or value within a JSON.

Other networks are specified in the same manner. The default file contains a public network, an admin network, a BMC vLAN, and a BMC network. Network configuration options are:

Table 16: Network Configuration Options

Name	Default	Description
mode	single	A string value of either single, dual, or team. This specifies the default network interface construction model.
teaming	map	A map of values specific to teaming
networks	map	A map of networks that this barclamp should manage



If you change the network bonding mode, manually reboot your nodes once the proposal has fully deployed. This makes the current active network configuration match the correctly created configuration files.

The system provides the following default networks:

Table 17: Default Networks

Name	Usage	Notes
admin	Private network for node-to-node communication	A router and router_pref is required; this network must be owned by the Crowbar system to run DHCP on it.
bmc	Private network for BMC communication	This can be made the same as the admin network by using the ranges to limit what IP goes where; a router, if wanted, is external to the system.
bmc_vlan	Private network for admin nodes on the BMC network	This must be the same as the BMC network and have the same vLAN. This will be used to generate a vLAN tagged interface on the Admin Nodes that can access the BMC LAN.
public	Public network for Crowbar and other components	Router and router_pref are required, router_pref should be the lowest.

Each network has the following parameters:

Table 18: Network Parameters

Name	Default	Description
vlan	Integer	This is the vLAN to use on the switch and interfaces for this network.
use_vlan	true	A value of "true" indicates that the vLAN should be applied to the interface. A value of "false" assumes that the node will receive untagged traffic for this network.
add_bridge	false	This indicates whether the network should have a bridge built on top of it. The bridge will be br.
subnet	IP Address	The subnet for this network
netmask	Netmask	The netmask for this network
router	IP Address	The default router for this network
router_pref	Integer	All networks with a gateway need a unique number. The lower the number will become the default route on the machine. This is not vlan-specific, but network-specific. For example, BMC

Name	Default	Description
		and BMC_VLAN should have different metric numbers.
broadcast	IP Address	The default broadcast address for this network
ranges	map	Contains a map of strings to start and stop values for the network; this allows for sub-ranges with the network for specific uses, e.g. dhcp, admin, BMC, hosts.

The range map has a string key that is the name and map defining the range.

Table 19: Range Map

Name	Type	Description
start	IP Address	First address in the range, inclusive
end	IP Address	Last address in the range, inclusive

JSON Configuration by Section

1. Attributes
 - a. Startup delay set to 30 seconds to allow spanning tree to settle down
 - b. Mode—This sets whether to build up single NICs or bonded NICs; options are single, teamed
 - c. Teaming—Sets the mode of the teaming: in this case 6
2. Interface Maps—To set up the interface map for figuring out and defining eth0, eth1, eth2 on particular hardware models:
 - a. Pattern—Pattern of the hardware model/type
 - b. Bus order—Order to start enumerating; enumeration begins at eth0, eth1, eth2, eth3. If a bus is not defined, it will be enumerated at the end in the order in which it was presented.
 - c. Conduit Maps—Determine what network gets mapped to which interface based on what role; pattern this matched to the attribute variable, the NIC type, and role
 - i. mode (single, team)
 - ii. NIC type 1g or 10g
 - iii. role—mastername, crowbar-config-default
3. Conduit Maps—determine how networks get mapped to which interface based on what role the first match will how networks are laid out.
 - a. Pattern- matches on three parts “Mode/NIC Type/Role Name”
 - i. mode single, team, wildcard
 - ii. NIC type and number in system. Example 1g2 matches the 2nd 1 gigabit port in the system, 10g3 matches the 3rd 10 gigabit port in the system
 - iii. role—This matches a role assigned to machine. For example.*namenode would match the name node role.
 - b. Conduit name (prod, mgmt, admin, intf0, intf1,...)
 - i. if_list — what interfaces to use
 - ii. team_mode — how to team when needed
 - c. repeat for other conduits
4. Networks — define the network, IP ranges, available scopes, etc.
 - a. Name of Network (admin, mgmt, prod)
 - i. conduit — Name of conduit used in step 4
 - ii. vlan — vLAN to use

- iii. use_vlan – weather to turn on 802.1q tagging
- iv. add_bridge – Whether to use bridging protocol or vLAN tagging
- v. subnet – The IP subnet
- vi. netmask – Subnet netmask
- vii. broadcast – Broadcast IP
- viii. router – default gateway for network
- ix. router_pref – likelihood of becoming the default route on a system
- x. ranges – The IP ranges in the subnet broken down by usage; admin, host, dhcp, are all possible examples



The following networks are required: bmc_lan, public, and admin. Admin must have ranges set to the dhcp, admin, and host.

How to Add a Public IP to a Node

From a command prompt on the Admin Node, you can execute the following:

- **crowbar network allocate_ip default <machine name> public host**

To validate address, you can run:

- **crowbar machines show <machine name>**

You should then have your system set up with a public IP. From the admin section above, you could do “admin switch” instead of “public host” and the IP allocated will be from the switch range of the admin network.

To edit the DNS or NTP time server, please modify the DNS and NTP barclamps.

Configuring the Network for Outbound Connectivity

The default configuration assumes the admin node is isolated, and does not depend on external connectivity. In some cases, it may be desirable to allow outbound access from the admin node, possibly to access external NTP or DNS servers. This requires updates to the public network definition. These changes need to be made before Crowbar is installed, and the IP needs to be allocated after Crowbar installation.

Updating the public network definition

Change the definition of the public network to match the external LAN before installing Crowbar. For example, in the following definition of the public network, lines 6-9, 12, and 13 would need to be updated to match the existing LAN. The entries that must be updated are the subnet, netmask, broadcast address, router, and ranges. You might also need to change the conduit entry, depending on which network interface is connected to the external network.

```

1      "public": {
2          "conduit": "intf1",
3          "vlan": 300,
4          "use_vlan": true,
5          "add_bridge": false,
6          "subnet": "192.168.126.0",
7          "netmask": "255.255.255.0",
8          "broadcast": "192.168.126.255",
9          "router": "192.168.126.1",
10         "router_pref": 5,
11         "ranges": {
12             "host": { "start": "192.168.126.2", "end": "192.168.126.49" },
13             "dhcp": { "start": "192.168.126.50", "end": "192.168.126.127" }
14         }
15     },

```

After saving the changes, finish the Crowbar installation.

Allocating the IP Address

After the Crowbar installation is complete, the last steps are to actually allocate the IP address and assign it to the interface. On the admin node, execute the commands:

```

# crowbar network allocate_ip default <admin name> public host
# chef-client
# /etc/init.d/chef-server-webui restart

```

Once the chef-client has finished, you should have access to the admin node through the new interface, and can outbound access to DNS or NTP servers.

If external access was being configured to support external DNS and/or NTP servers, then edit any new or existing barclamp proposals to include the external server entries, and apply the proposals.

To verify NTP access, you can use the ntpq utility. The '*' before the node IP indicates the local NTP client has synchronized to the external server. You should wait for the admin node to initially sync to the external time server before deploying the remaining nodes.

```

[root@admin config]# ntpq -p
remote refid st t when poll reach delay offset jitter
=====
*172.26.1.50 132.163.4.103 2 u 40 64 377 0.287 -0.433 0.169

```


Verifying Master Node State

When the Admin Node finishes installation, it will remain at a shell prompt. At this point, all Crowbar services have started. Consult the table below to access these services.

Table 20: Crowbar Services

Service	URL	Credentials
SSH	root@192.168.124.10	crowbar
Crowbar UI	http://192.168.124.10:3000/	crowbar / crowbar
Nagios	http://192.168.124.10/nagios3	nagiosadmin / password
Ganglia	http://192.168.124.10/ganglia	nagiosadmin / password
Chef UI	http://192.168.124.10:4040/	admin / password

Logging into the UI requires acceptance of the EULA. It can be found on the dashboard under EULA, and in Appendix A of this document.

Data Node Installation

Nodes other than the Admin Nodes are installed when they are first powered up. A sequence boot phase is executed (rebooting multiple times) which culminates in deploying a minimal OS image installed on the local drive. Part of the basic installation includes “hooking” the nodes into the infrastructure services—NTP, DNS, Nagios, and Ganglia.

Once known to Crowbar, the node can be managed; it can be powered on and off, rebooted, and components can be installed on it.

Functional components are installed on nodes by including them in one or more barclamps’ proposals. For example, when a proposal names a Master name node, the relevant packages, services, and configuration are deployed to that node when the proposal is committed.

Troubleshooting a Hardware Installation Loop

If you find that some nodes are in a hardware installation loop, with the UI status continuously changing between *Problem* and *Hardware Installing*, the following procedure can help resolve the issue.

1. Run `knife node list`.
2. Set your editor (`export EDITOR=<editor_of_choice>`).
3. Run `knife node edit <nodename>`.
4. Find the setting "track": {

```

"bios_c_wsman_attempts": 3,
"bios_wsman_attempts": 3,
"bios_wsman_u_attempts": 3

```
5. Change those values from **3** to **0**, and then close the editor.
6. Reboot the node.

The nodes should auto-correct after the next state transition to *Hardware Installing*.

The next section describes details for installing the different components.

Installing components

The general workflow to install any component is the same:

1. Obtain a default proposal that includes the parameters for the component and a mapping of nodes to the roles they are assigned.
2. Edit the proposal to match the desired configuration.
3. Upload the proposal to Crowbar.
4. Commit the proposal.
5. Reboot the nodes.

All these activities are achieved by using the Crowbar command line tool or the Web-based UI. The sections that follow use the command line tool: `/opt/dell/bin/crowbar`.

In the sections that follow, this tool is referred to as "Crowbar."

General Installation Process

Set CROWBAR Parameter

Before any of the following command lines may be used, the CROWBAR_KEY environment variable must be primed with the SSH key so they may connect.

```
On the admin node, use this command: export CROWBAR_KEY=crowbar:crowbar
On all other nodes, use this: export CROWBAR_KEY=$(cat /etc/crowbar.install.key)
```

Obtain a proposal

Crowbar can inspect the current known nodes and provide a proposal that best utilizes the available systems for the component being installed. To obtain and inspect this proposed configuration:

```
/opt/dell/bin/crowbar <component> proposal create <name>
/opt/dell/bin/crowbar <area> proposal show <name> > <local_file_name>
```

Where:

- **<area>** — The area for which the proposal is made (e.g., Clouderamanager, Pig).
- **<name>** — The name assigned to this proposal. This name should be unique for the component; i.e. if two hadoop clusters are being installed, the proposals for each should have unique names.
- **<local_file_name>** — Any file name into which the proposal will be written.

Update a proposal

The local file created above can be inspected and modified. The most common changes are:

- **Change default passwords and other barclamp parameters (e.g. swift replica count).**
- **Change assignment of machines to roles.**

Once edits are completed, Crowbar must be updated. To update Crowbar with a modified proposal, execute:

```
/opt/dell/bin/crowbar <area> proposal --file=<local_file_name> edit <name>
```

Where the parameters in this command are exactly as mentioned above, Crowbar will validate the proposal for syntax and basic sanity rules as part of this process.

Committing a proposal

Once the proposal content is satisfactory, the barclamp instance can be activated. To achieve that, execute:

```
/opt/dell/bin/crowbar <area> proposal commit <name>
```

This might take a few moments, as Crowbar is deploying the required software to the machines mentioned in the proposal.

Modifying an active configuration

When committing a proposal that was previously committed, Crowbar compares the new configuration to the currently active state and applies the deltas.

To force Crowbar to reapply a proposal, the active state needs to be deleted:

```
/opt/dell/bin/crowbar <area> delete <name>
```

Installing Cloudera Manager

Cloudera Manager is installed as part of the Clouderamanager barclamp. To install it, refer to the *Dell | Cloudera Solution User's Guide*.

Hadoop Ecosystem Components

Some ecosystem components are installed via Cloudera Manager (HUE, Oozie, Hbase, Hive, Sqoop, and Zookeeper). The *Dell | Cloudera Solution User's Guide* shows how these services are installed or added to a Hadoop cluster. Pig must be installed via Crowbar barclamps.

Table 21: Hadoop Ecosystem Components

Component	Master Node	Data Node	Edge Node	Utilize From	Administer From
Pig	X	X	X	Edge Node	Edge Node
Hive		X	X	Edge Node	Edge Node
Sqoop			X	Edge Node	Edge Node



"X" designates server location for the appropriate package binaries to be installed.

Dell | Cloudera Solution Monitoring and Alerting

Table 22: Components Monitored by Hadoop Monitoring Console

Service Type	Resource	Warning	Critical	Nodes to Monitor
Disk	HDFS_DISK_[00-10]	60	90	DataNode[]
SWAP	SWAP	60	90	DataNode[]
		60	90	Master Node[]
		60	90	EdgeNode[]
Ping_Node_From_Admin		DELAY	NO RESPONSE	DataNode[]
		DELAY	NO RESPONSE	Master Node[]
		DELAY	NO RESPONSE	EdgeNode[]
NIC Bonding		DELAY	1 NIC in Bond	DataNode[]
		DELAY	1 NIC in Bond	Master Node[]
		DELAY	1 NIC in Bond	EdgeNode[]
DNS_From_Node		DELAY	NO RESPONSE	DataNode[]
		DELAY	NO RESPONSE	Master Node[]
		DELAY	NO RESPONSE	EdgeNode[]
DNS_About_Node		DELAY	NO RESPONSE	DataNode[]
		DELAY	NO RESPONSE	Master Node[]
		DELAY	NO RESPONSE	DataNode[]
JobTracker_Daemon		DELAY	DAEMON NOT RUNNING	Master Node[]
TaskTracker_Daemon		DELAY	DAEMON NOT RUNNING	DataNode[]
DataNode_Daemon		DELAY	DAEMON NOT RUNNING	DataNode[]
Master Node_Daemon		DELAY	DAEMON NOT RUNNING	Master Node[]
SecondaryMaster Node		DELAY	DAEMON NOT RUNNING	Master Node[]
SSH		DELAY	NO RESPONSE	DataNode[]
		DELAY	NO RESPONSE	Master Node[]

Service Type	Resource	Warning	Critical	Nodes to Monitor
Zombie_Processes		5	10	DataNode[]
		5	10	Master Node[]
		5	10	EdgeNode[]
CPU_Load		80	90	DataNode[]
		80	90	Master Node[]
		80	90	EdgeNode[]
Zookeeper_Client		DELAY	DAEMON NOT RUNNING	DataNode[]
Zookeeper_Server		DELAY	DAEMON NOT RUNNING	Master Node[]
JobTracker_Submit_Job		DELAY	NO RESPONSE	Master Node[]
Chef_Daemon		DELAY	NO RESPONSE	DataNode[]
		DELAY	NO RESPONSE	Master Node[]
		DELAY	NO RESPONSE	EdgeNode[]
Disk	MAPRED_DIR	60	90	DataNode[]
		60	90	Master Node[]
		60	90	EdgeNode[]
Memory_Capacity_Used	System Memory	80	90	DataNode[]
		80	90	Master Node[]
		80	90	EdgeNode[]
Disk	HDFS01_Capacity	60	90	Master Node[]
				DataNode[]
				Master Node[]
CPU_Utilization				DataNode[]
				Master Node[]
				EdgeNode[]
Memory_Utilization				DataNode[]
				Master Node[]
				EdgeNode[]
NIG_LAG_Utilization				DataNode[]
				Master Node[]

Service Type	Resource	Warning	Critical	Nodes to Monitor
				EdgeNode[]
CPU Temp		As defined by SDR (Sensor Data Record)	As defined by SDR	DataNode[]
		As defined by SDR	PENDING	Master Node[]
		As defined by SDR	As defined by SDR	EdgeNode[]
Power Supplies		As defined by SDR	As defined by SDR	Master Node[]
		As defined by SDR	As defined by SDR	Edge Node[]
Master Node _NFS_Mount		DELAY	MOUNT MISSING	Master Node[]
Hbase		DELAY	SELECT FAILED	EdgeNode[]
		DELAY	INSERT FAILED	EdgeNode[]
Hive		DELAY	SELECT FAILED	EdgeNode[]
		DELAY	INSERT FAILED	EdgeNode[]
Ping_From_Admin	IPMI Interface	DELAY	NO RESPONSE	DataNode[]
		DELAY	NO RESPONSE	Master Node[]
		DELAY	NO RESPONSE	EdgeNode[]

Appendix A: The Dell End User License Agreement

DELL END USER LICENSE AGREEMENT Ver. 1.3 TYPE-A SOFTWARE

THIS IS A LEGAL AGREEMENT BETWEEN YOU (EITHER AN INDIVIDUAL OR AN ENTITY) AND DELL PRODUCTS L.P. OR DELL GLOBAL B.V. (SINGAPORE BRANCH), ON BEHALF OF DELL INC. AND ITS WORLDWIDE SUBSIDIARIES AND AFFILIATES (COLLECTIVELY, "Dell" OR "DELL"), WHICH GOVERNS YOUR USE OF THE SOFTWARE. THE SOFTWARE SHALL MEAN COLLECTIVELY THE SOFTWARE PROGRAM, THE ASSOCIATED MEDIA, PRINTED MATERIALS, ONLINE OR ELECTRONIC DOCUMENTATION, AND ANY COPIES THEREOF, TO WHICH THIS AGREEMENT IS ATTACHED OR OTHERWISE ASSOCIATED (the "Software" or "SOFTWARE"). PLEASE READ THE TERMS AND CONDITIONS OF THIS AGREEMENT CAREFULLY, INCLUDING, WITHOUT LIMITATION, ANY SUPPLEMENTAL TERMS AND CONDITIONS APPEARING OR REFERENCED BELOW, WHICH ARE HEREBY MADE PART OF THIS END USER LICENSE AGREEMENT (COLLECTIVELY, "EULA"), BEFORE DOWNLOADING, INSTALLING, ACTIVATING AND/OR OTHERWISE USING THE SOFTWARE. BY EXPRESSLY ACCEPTING THESE TERMS OR DOWNLOADING, INSTALLING, ACTIVATING AND/OR OTHERWISE USING THE SOFTWARE, YOU ARE AGREEING THAT YOU HAVE READ, AND THAT YOU AGREE TO COMPLY WITH AND TO BE BOUND BY THE TERMS AND CONDITIONS OF THIS EULA AND ALL APPLICABLE LAWS AND REGULATIONS IN THEIR ENTIRETY WITHOUT LIMITATION OR QUALIFICATION. IF YOU DO NOT AGREE TO BE BOUND BY THE TERMS AND CONDITIONS OF THIS EULA, THEN YOU MAY NOT DOWNLOAD, INSTALL, ACTIVATE OR OTHERWISE USE ANY OF THE SOFTWARE AND YOU MUST PROMPTLY RETURN THE SOFTWARE AND ANY HARDWARE TO WHICH IT IS ATTACHED, AS DIRECTED BY DELL OR ITS RESELLER (IF APPLICABLE) FOR A FULL REFUND. IF YOU ARE AN INDIVIDUAL REPRESENTING AN ENTITY, YOU ACKNOWLEDGE THAT YOU HAVE THE APPROPRIATE AUTHORITY TO ACCEPT THESE TERMS AND CONDITIONS ON BEHALF OF SUCH ENTITY.

1. **License** Subject to the terms, conditions and limitations of this EULA and timely payment by you, Dell hereby grants you a limited, nonexclusive, nontransferable, non-assignable license, without rights to sublicense, to install or have installed, display and use the Software (in object code only) only on as many computers, devices and/or in such configurations as you are expressly entitled (e.g., as set forth in the applicable Dell sales quote). The terms and conditions of this EULA will govern use of the Software and any upgrades, updates, patches, hotfixes, modules, routines and/or additional versions of the Software provided by Dell, **at Dell's sole discretion, that replace and/or supplement the original Software (collectively, "Update"), unless such Update is accompanied by or references a separate license agreement in which case the terms and conditions of that agreement will govern.** If this EULA governs your use of an Update, such Update shall be considered Software for purposes of this EULA. Unless earlier terminated as provided herein, the term of each individual license granted under this EULA begins on the date of acceptance of this EULA, and continues only for such period as you have purchased, in the case of a term license, and is perpetual if no term is specified. Customer may use only portions of Software for which it has paid the applicable license fee.
2. **License Limitations** You may not copy the Software except for a reasonable number of copies solely **as needed for backup or archival purposes or as otherwise expressly permitted in Section 1 "License"** above. You may not modify or remove any titles, trademarks or trade names, copyright notices, legends, or other proprietary notices or markings on or in the Software. The rights granted herein are **limited to Dell's and its licensors' and suppliers' intellectual property rights in the Software and do not include any other third party's intellectual property rights. If the Software was provided to you on removable media (e.g., CD, DVD, or USB drive), you may own the media on which the Software is recorded but Dell, Dell's licensor(s) and/or supplier(s) retain ownership of the Software itself and all related intellectual property rights.** You are not granted any rights to any trademarks or service marks of Dell. This EULA does not apply to any third party software that is not included as part of the Software. The use of any other software, including any software package or file, whether licensed to you separately by Dell or by a third party, is subject to the terms and conditions that come with or are associated with such software.

3. **Rights Reserved.** THE SOFTWARE IS LICENSED, NOT SOLD. Except for the license expressly granted in this EULA, Dell, on behalf of itself and its licensors and suppliers, retains all right, title, and interest in and to the Software and in all related content, materials, copyrights, trade secrets, patents, trademarks, derivative works and any other intellectual and industrial property and proprietary rights, including registrations, applications, renewals, and extensions of such rights (the "Works"). The rights in these Works are valid and protected in all forms, media and technologies existing now or hereafter developed and any use other than as contemplated herein, including the reproduction, modification, distribution, transmission, adaptations, translation, display, republication or performance of the Works, except as specifically permitted herein, is strictly prohibited. Dell, on behalf of itself and its licensors and suppliers, retains all rights not expressly granted herein.
4. **Restrictions.** Except as otherwise provided herein or expressly agreed by Dell, you may not, and will not allow a third party to (A) sell, lease, license, sublicense, assign, distribute or otherwise transfer or encumber in whole or in part the Software; (B) provide, make available to, or permit use of the Software in whole or in part by, any third party, including contractors, without Dell's prior written **consent, unless such use by the third party is solely on Customer's behalf, is strictly in compliance** with the terms and conditions of this EULA, and you are liable for any breach of this EULA by such **third party (a "Permitted Third Party")**; (C) **copy, reproduce, republish, upload, post or transmit the** Software in any way (D) decompile, disassemble, reverse engineer, or otherwise attempt to derive source code (or underlying ideas, algorithms, structure or organization) from the Software program, in whole or in part; (E) modify or create derivative works based upon the Software (except as set forth in Section 13 – Development Tools); (F) use the Software on a service bureau, rental or managed services basis or permit other individuals or entities to create Internet "links" to the Software or "frame" or "mirror" the Software on any other server or wireless or Internet-based device; (G) use Software that was loaded by Dell onto specific hardware (an "Appliance") **separately from such** Appliance; or (H) use the Software to create a competitive offering. You may not, and will not allow a Permitted Third Party to, use the Software in excess of the number of licenses expressly authorized by Dell. **In addition, you may not share the results of any benchmarking of the Software without Dell's** prior written consent.
5. **Compliance.** You will certify in writing, upon reasonable request by Dell, that all use of Software is in compliance with the terms of this EULA, indicating the number of Software licenses deployed at that time. You grant Dell, or an agent selected by Dell, the right to perform a reasonable audit of your compliance with this EULA during normal business hours. You agree to cooperate and provide Dell with all records reasonably related to your compliance with this EULA. If, as a result of the audit, a deficiency of greater than five percent (5%) is found in the licensee fees paid, then you shall bear the total cost of the audit, in addition to any other liabilities you may have.
6. **Support and Subscription Services Not Included.** Dell does not provide any maintenance or support services under this EULA. Maintenance and support services, if any, are provided under a separate agreement, which may be located at www.Dell.com/servicecontracts. Additionally, this EULA, in and of itself, does not entitle you to any Updates at any time in the future.
7. **Termination.** Dell may terminate this EULA immediately and without prior notice if you fail to comply with any term or condition of this EULA or if Dell does not receive timely payment for the licenses to the Software or for the hardware to which it is attached, if any. In addition, Dell may terminate any license to Software distributed for free at any time in its sole discretion. Either party may terminate this EULA at any time by providing at least ninety (90) days prior written notice to the other party. In the event of termination of this EULA, all licenses granted hereunder shall automatically terminate and you must immediately cease use of the Software and return or destroy all copies of the Software. The parties recognize and agree that their obligations under Sections 3, 4, 7, 10, 11, 12, 17, 18, 19, 20, 22 and 23 of this EULA, as well as obligations for payment, survive the cancellation, termination, and/or expiration of this EULA, and/or the licenses granted hereunder.

8. **Export, Import and Government Restrictions.** You are advised that the Software is subject to U.S. export laws as well as the laws of the country where it is delivered or used. You agree to abide by these laws. Under these laws, the Software may not be sold, leased, or transferred to restricted countries (currently Cuba, Iran, North Korea, Sudan and Syria), restricted end-users, or for restricted end-uses. You specifically agree that the Software will not be used for activities related to weapons of mass destruction, including but not limited to, activities related to the design, development, production or use of nuclear materials, nuclear facilities, or nuclear weapons, missiles or support of missile projects, or chemical or biological weapons. You understand that certain functionality of the Software, such as encryption or authentication, may be subject to import restrictions in the event you transport the Software from the country of delivery and you are responsible for complying with applicable restrictions.

The Software and documentation are "commercial items" as that term is defined at 48 C.F.R. 2.101, consisting of "commercial computer software" and "commercial computer software documentation" as such terms are used in 48 C.F.R. 12.212. Consistent with 48 C.F.R. 12.212 and 48 C.F.R. 227.7202-1 through 227.7202-4, all U.S. Government end users acquire the Software and documentation with only those rights set forth herein. Contractor/manufacture is Dell Products L.P., One Dell Way, Round Rock, Texas 78682.

9. **Limited Warranty.** Unless otherwise provided in the limited warranties for the Software found at <http://www.Dell.com/Warranty>, Dell solely warrants that it has the right to grant the licenses to the Software, and except as set forth in Sections 14 and 16 below, that such Software will substantially conform in material respects to the functional specifications and current documentation provided by Dell with the Software. This limited warranty is not transferable and extends only for thirty (30) days from the date of delivery of the Software, unless otherwise stated at www.Dell.com/Warranty. This limited warranty does not cover damages, defects, malfunctions or failures caused by any unauthorized modification by you, or your agents, of the Software; any abuse, misuse or negligent acts of you; modification by you of any interfaces or any software or hardware interfacing with the Software; or **any failure by you to follow Dell's installation, operation or maintenance instructions.** EXCEPT FOR THE PRECEDING EXPRESS LIMITED WARRANTY, DELL MAKES, AND YOU RECEIVE, NO OTHER WARRANTIES RELATED TO THE SOFTWARE WHETHER EXPRESS, IMPLIED OR STATUTORY, AND DELL SPECIFICALLY DISCLAIMS ANY IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, AND NON-INFRINGEMENT. DELL DOES NOT WARRANT THAT THE FUNCTIONS OF THE SOFTWARE WILL MEET YOUR REQUIREMENTS OR THAT OPERATION OF THE SOFTWARE WILL BE UNINTERRUPTED OR ERROR FREE. YOU ASSUME RESPONSIBILITY FOR SELECTING THE SOFTWARE AND THE RESULTS ACHIEVED. YOUR SOLE AND EXCLUSIVE REMEDY, **AND DELL'S ENTIRE LIABILITY, FOR BREACH OF THE WARRANTIES PROVIDED HEREIN, IS FOR DELL,** AT ITS SOLE DISCRETION, TO EITHER USE COMMERCIALY REASONABLE EFFORTS TO REMEDY ANY NON-CONFORMANCE OR TO PROVIDE A REFUND OF THE LICENSE FEES PAID BY YOU TO DELL FOR THE SOFTWARE. THIS DISCLAIMER OF WARRANTY MAY NOT BE VALID IN SOME JURISDICTIONS AND YOU MAY HAVE WARRANTY RIGHTS UNDER LAW WHICH MAY NOT BE WAIVED OR DISCLAIMED -- ANY SUCH WARRANTY EXTENDS ONLY FOR THIRTY (30) DAYS FROM THE DATE OF DELIVERY OF THE SOFTWARE.
10. **Limitation of Liability.** DELL WILL NOT BE LIABLE FOR ANY INCIDENTAL, INDIRECT, PUNITIVE, SPECIAL OR CONSEQUENTIAL DAMAGES, ARISING OUT OF OR IN CONNECTION WITH THIS EULA AND/OR THE SOFTWARE. DELL SHALL HAVE NO LIABILITY FOR THE FOLLOWING: (A) LOSS OF REVENUE, INCOME, PROFIT, OR SAVINGS, (B) LOST OR CORRUPTED DATA OR SOFTWARE, LOSS OF USE OF SYSTEM(S) OR NETWORK(S), OR THE RECOVERY OF SUCH, (C) LOSS OF BUSINESS OPPORTUNITY, (D) BUSINESS INTERRUPTION OR DOWNTIME, (E) LOSS OF GOODWILL OR REPUTATION, OR (F) SOFTWARE NOT BEING AVAILABLE FOR USE OR THE PROCUREMENT OF SUBSTITUTE SOFTWARE OR GOODS.

NOTWITHSTANDING ANYTHING TO THE CONTRARY SET FORTH IN THIS EULA, DELL'S TOTAL LIABILITY FOR ANY AND ALL CLAIMS ARISING OUT OF OR IN CONNECTION WITH THIS EULA AND/OR THE SOFTWARE SHALL NOT EXCEED THE TOTAL AMOUNT RECEIVED BY DELL FOR THE PARTICULAR SOFTWARE GIVING RISE TO SUCH CLAIM(S).

The foregoing limitations, exclusions and disclaimers shall apply to any and all claims, regardless of whether the claim(s) for such damages is based in contract, warranty, strict liability, negligence, tort or otherwise. Insofar as applicable law prohibits any limitation on liability herein, the parties agree that such limitation will be automatically modified, but only to the extent so as to make the limitation compliant with applicable law. The parties agree that the limitations on liabilities set forth herein are agreed allocations of risk and such limitations will apply notwithstanding the failure of essential purpose of any limited remedy and even if a party has been advised of the possibility of any such liability.

DELL SHALL NOT BE LIABLE TO YOU FOR ANY CLAIM BROUGHT MORE THAN TWO YEARS AFTER THE CAUSE OF ACTION FOR SUCH CLAIM FIRST AROSE.

11. **Indemnification**. Except where prohibited by applicable law, Dell shall defend and indemnify you against any third-party claim or action that the Software (specifically excluding open source software) **infringes or misappropriates that third party's patent, copyright, trade secret, or other intellectual property rights ("Indemnified Claims")**. In addition, if Dell receives prompt notice of an Indemnified Claim that, in Dell's reasonable opinion, is likely to result in an adverse ruling, then Dell shall at its sole discretion, (A) obtain a right for you to continue using such Software; (B) modify such Software to be non-infringing; (C) replace such Software with a non-infringing substitute; or (D) provide a reasonable depreciated refund for the allegedly infringing Software. Notwithstanding the foregoing, Dell shall have no obligation under this Section for Indemnified Claims resulting or arising from: (i) modifications of the Software that were not performed by or on behalf of Dell; (ii) your unauthorized use of the Software, or the combination or operation of the Software in connection with a third-party product, software, or service (the combination of which causes the claimed infringement); (iii) your failure to incorporate free Software updates or upgrades that would have avoided the alleged infringement; or (iv) Dell's compliance with your written specifications or directions, including the incorporation of any software or other materials or processes provided by or requested by you (collectively, "Excluded Indemnified Claims"). Dell's duty to indemnify and defend is contingent upon: (a) your providing Dell with prompt written notice of the third-party claim or action, (b) Dell having the right to solely control the defense and settlement of such claim or action, and (c) your cooperation with Dell in defending and resolving such claim or action. This Section states your exclusive remedies for any third-party claim or action, and nothing in this EULA or elsewhere will obligate Dell to provide any greater indemnity to you. You, at your expense, shall defend and indemnify Dell against any claim, action or proceeding brought against Dell which arises from or is in any manner connected with the Excluded Indemnified Claims.
12. **Confidentiality**. You agree to: (A) not use Confidential Information except as necessary to exercise the rights herein and (B) use best efforts to preserve and protect the confidentiality of the Confidential Information. "Confidential Information" means any oral, written, graphic or machine-readable information disclosed by Dell that is (i) identified as confidential; (ii) designated in writing to be confidential or proprietary; or (iii) should be reasonably understood to be confidential. Confidential Information includes the Software and its trade secrets, including but not limited to source code, the development status of the Software, the appearance, content and flow of the user interface of the Software, and the content and documentation of the Software. Confidential Information does not include information that is (a) publicly available other than through a breach of this EULA; (b) known to you prior to such disclosure; or (c) subsequently lawfully obtained by you from a third party that **has no obligations of confidentiality. You agree that, without Dell's prior written consent, you will not grant access to any Dell Confidential Information to any persons or entities except for your employees and agents who have a business need to have such access and who are obligated to maintain the**

confidentiality thereof as set forth herein. These obligations do not expire. In some, limited circumstances, Dell may need to engage a third party to fulfill its obligations to you under this EULA. By using this Software you agree that Dell may provide your information to such third party for that purpose. Dell may obtain information related to your use of the Software and you agree that we may use such information in aggregate form in an anonymous manner in support of our marketing activities related to the Software. Any feedback, value added changes or suggestions made by you or other information that is provided to Dell relating to the Software shall be owned by Dell and considered Dell Confidential Information.

13. **Development Tools**. If the Software includes development tools, such as scripting tools, APIs, or sample scripts (collectively "Development Tools"), you may use such Development Tools to create new scripts and code for the purpose of customizing your use of the Software (within the parameters set forth in this EULA and within the parameters set forth in the Development Tools themselves) and for no other purpose. Notwithstanding anything to the contrary set forth in this EULA, no warranty or technical support is provided for sample scripts contained in such Development Tools or scripts or other code written by you or any third party.
14. **Evaluation Licenses**. If you have received Software for evaluation purposes ("Evaluation Software"), you may use the Evaluation Software solely for such limited evaluation period and for internal evaluation purposes only. You acknowledge that Dell may terminate your right to evaluate the Evaluation Software, for any or no reason, effective immediately upon notice to you. IN ADDITION, THE EVALUATION SOFTWARE IS PROVIDED TO YOU "AS IS" WITHOUT INDEMNITY OR WARRANTY OF ANY KIND, WHETHER EXPRESS, IMPLIED, STATUTORY, OR OTHERWISE. NOTWITHSTANDING ANYTHING TO THE CONTRARY IN THIS EULA, DELL BEARS NO LIABILITY FOR ANY DIRECT, INDIRECT, INCIDENTAL, PUNITIVE, SPECIAL OR CONSEQUENTIAL DAMAGES RESULTING FROM USE (OR ATTEMPTED USE) OF THE EVALUATION SOFTWARE THROUGH AND AFTER THE EVALUATION PERIOD AND HAS NO DUTY TO PROVIDE SUPPORT TO YOU.
15. **Hosted and Internet-Accessible Software**. Some or all of the Software may be remotely hosted or accessible to you through the Internet ("Hosted Software"). In such case, Dell may suspend, terminate, withdraw, or discontinue all or part of the Hosted Software or your access to the Hosted Software upon receipt of a subpoena or law-enforcement request, or when Dell believes, in its sole discretion, that you have breached any term of this EULA or are involved in any fraudulent, misleading, or illegal activities. Dell may modify the Hosted Software at any time with or without prior notice to you. Dell may perform scheduled or unscheduled repairs or maintenance, or remotely patch or upgrade the Hosted Software installed on its and your system(s), which may temporarily degrade the quality of the Hosted Software or result in a partial or complete outage of the Hosted Software. Updates, patches or alerts may be delivered from Dell servers, which may be located outside of your country. Dell provides no assurance that you will receive advance notification of such activities or that your use of the Hosted Software will be uninterrupted or error-free.
16. **Open Source Software**. The Software may come bundled or otherwise be distributed with open source software, which is subject to the terms and conditions of the specific license under which the open source software is distributed. THIS OPEN SOURCE SOFTWARE IS PROVIDED BY DELL "AS IS" WITHOUT ANY WARRANTY, EXPRESS, IMPLIED, OR OTHERWISE, INCLUDING BUT NOT LIMITED TO THE IMPLIED WARRANTY OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NON-INFRINGEMENT. NOTWITHSTANDING ANYTHING TO THE CONTRARY IN THIS EULA, AS IT RELATES TO ANY AND ALL CLAIMS ARISING OUT OF OR IN CONNECTION WITH OPEN SOURCE SOFTWARE, DELL SHALL HAVE NO LIABILITY FOR ANY DIRECT, INDIRECT, INCIDENTAL, PUNITIVE, SPECIAL OR CONSEQUENTIAL DAMAGES. Under certain open source software licenses, you are entitled to obtain the corresponding source files. You may find corresponding source files for the Software at <http://opensource.dell.com> or other locations that may be specified to you by Dell.
17. **Right to Preliminary and Injunctive Relief**. You agree that money damages would be an inadequate remedy for Dell in the event of a breach or threatened breach by you of the provisions set

forth in this EULA; therefore, you agree that in the event of a breach or threatened breach of any such provisions, Dell may, in addition to any other remedies to which it is entitled, be entitled to such preliminary or injunctive relief (including an order prohibiting you from taking actions in breach of **such provisions**) and **specific performance as may be appropriate to preserve all of Dell's rights**. All rights and remedies afforded Dell by law shall be cumulative and not exclusive.

18. **Choice of Law and Language**. This EULA shall be governed by the laws of the State of Texas, USA, to the exclusion of the UN Convention on Contracts for the International Sale of Goods. You acknowledge that the headquarters of the Dell family of companies is located in Texas, and that the software licensed under this EULA and the related products marketed in connection with such software were in substantial part conceived, developed, and marketed by Dell personnel in Texas. Further, you acknowledge, agree, and stipulate that the laws of the State of Texas bear a substantial relationship to this EULA and that the selection of Texas law to govern this EULA and the license of the Software hereunder is reasonable and appropriate, and you consent to the selection of such law to govern this EULA and the relationship of the parties hereto. This EULA has been agreed only in the English language, which version of this EULA shall be controlling regardless of whether any translations of this EULA have been prepared or exchanged. As an exception to the preceding sentence, if Dell provides this EULA to you only in a non-English language version, then such non-English language version shall control. You acknowledge and represent that you have carefully reviewed this EULA with the involvement and assistance of your employees, advisors, and/or legal counsel fluent in the English language, that you have consulted with local legal counsel and counsel competent to render advice with respect to transactions governed by the law applicable to this EULA, that you have no questions **regarding the meaning or effect of any of this EULA's terms, and that you have obtained high-quality translations of this EULA for use by you or any of your team who are not fluent in the English language, with the understanding that you alone shall bear the risk of any misunderstandings that may arise as a result of such translation.** All communications in connection with this EULA shall be in the English language.

Les parties ont demandé que cette convention ainsi que tous les documents qui s'y rattachent soient rédigés en anglais.

19. **Dispute Resolution and Binding Arbitration**. ANY CLAIM, DISPUTE, OR CONTROVERSY (WHETHER IN CONTRACT, TORT, OR OTHERWISE, WHETHER PREEXISTING, PRESENT OR FUTURE, AND INCLUDING STATUTORY, COMMON LAW, INTENTIONAL TORT AND EQUITABLE CLAIMS) BETWEEN YOU AND DELL arising out of or in connection with this EULA, or the breach, termination or validity thereof shall be finally settled under the Rules of Arbitration of the International Chamber of Commerce ("ICC") by one or more arbitrators with expertise in software licensing appointed in accordance with such rules. The arbitration shall be conducted in the English language. The place of the arbitration shall be a commercial center reasonably chosen by the arbitration panel in a third country so as to ensure that the award resulting from the arbitration shall be of an international character and enforceable under the New York Convention on the Recognition and Enforcement of Foreign Arbitral Awards. The arbitration panel shall be empowered to grant whatever relief would be available in court, including without limitation preliminary relief, injunctive relief, and specific performance. Any award of the arbitration panel shall be final and binding immediately when rendered, and judgment on the award may be entered in any court of competent jurisdiction. Neither you nor Dell shall be entitled to join, consolidate, or include any claims belonging to or alleged or arising from, by, or on behalf of any third party to an arbitration brought hereunder. The individual (non-class) nature of this dispute resolution provision goes to the essence of the parties' dispute resolution agreement, and if found unenforceable, the entire arbitration and dispute resolution provision shall be void. Notwithstanding the foregoing, Dell may apply to any relevant government agency or any court of competent jurisdiction to preserve its rights under this EULA and to obtain any injunctive or preliminary relief, or any award of specific performance, to which it may be entitled, either against you or against a non-party; provided, however, that no such administrative or judicial authority shall have the right or power to render a judgment or award (or to enjoin the rendering of

an arbitral award) for damages that may be due to or from either party under this EULA, which right and power shall be reserved exclusively to an arbitration panel proceeding in accordance herewith.

20. **No Waiver**. No waiver of breach or failure to exercise any option, right, or privilege under the terms of this EULA on any occasion shall be construed to be a waiver of a subsequent breach or right to exercise any option, right, or privilege.
21. **Force Majeure**. Dell shall not be responsible for any delay or failure in performance of any part of this EULA to the extent that such delay or failure is caused by fire, flood, explosion, war, embargo, government requirement, civil, or military authority, act of God, act or omission of carriers, failure of the Internet or other similar causes beyond its control.
22. **No Assignment**. Except as set forth herein, you may not assign or transfer your interests, rights or obligations under this EULA by written agreement, merger, consolidation, operation of law or otherwise, without the prior written consent of an authorized executive officer of Dell. Any attempt to assign this EULA by you without such prior written consent from Dell shall be null and void.
23. **Entire Agreement**. Unless you have entered into another written agreement with respect to the Software which has been signed by you and an authorized representative of Dell and which conflicts with the terms of this EULA, you agree that this EULA supersedes all prior written or oral agreements, warranties or representations, including any and all other click-wrap, shrink-wrap or similar licenses or agreements, with respect to the Software. No amendment to or modification of this EULA, in whole or in part, will be valid or binding unless it is in writing and executed by authorized representatives of both parties. If any term of this EULA is found to be invalid or unenforceable, the remaining provisions will remain effective. You agree that any principle of construction or rule of law that provides that an agreement shall be construed against the drafter shall not apply to the terms and conditions of this EULA.
24. **Notices**. Notice to Dell under this EULA must be in writing and sent to the address below or to such other address (including facsimile or e-mail) as specified in writing, and will be effective upon receipt.

**Dell Inc., Attn: Dell Legal
One Dell Way, Round Rock, Texas 78682**

Last rev. 021812

Supplemental Terms and Conditions:

Appendix B: Network JSON Example

```
{
  "id": "bc-template-network",
  "description": "Instantiates network interfaces on the crowbar managed systems. Also
manages the address pool",
  "attributes": {
    "network": {
      "start_up_delay": 30,
      "mode": "team",
      "teaming": {
        "mode": 6
      }
    },
    "interface_map": [
      {
        "pattern": "PowerEdge R610",
        "bus_order": [
          "0000:00/0000:00:01",
          "0000:00/0000:00:03"
        ]
      },
      {
        "pattern": "PowerEdge R710",
        "bus_order": [
          "0000:00/0000:00:01",
          "0000:00/0000:00:03"
        ]
      },
      {
        "pattern": "PowerEdge C6145",
        "bus_order": [
          "0000:00/0000:00:04",
          "0000:00/0000:00:02"
        ]
      },
      {
        "pattern": "PowerEdge C2100",
        "bus_order": [
          "0000:00/0000:00:1c",
          "0000:00/0000:00:07",
          "0000:00/0000:00:09",
          "0000:00/0000:00:01"
        ]
      },
      {
        "pattern": "C6100",
        "bus_order": [
          "0000:00/0000:00:01",
          "0000:00/0000:00:03",
          "0000:00/0000:00:07"
        ]
      },
      {
        "pattern": "product",
        "bus_order": [
          "0000:00/0000:00:01",
          "0000:00/0000:00:02"
        ]
      }
    ],
    "conduit_map": [
      {
        "pattern": "team./*/crowbar-config-default",
        "conduit_list": {
          "prod": {
```

```

        "if_list": [ "10g1", "10g3" ],
        "team_mode": 6
    },
    "public": {
        "if_list": [ "10g2", "10g4" ],
        "team_mode": 6
    },
    "mgmt": {
        "if_list": [ "1g1" ]
    }
},
{
    "pattern": "team/*/*.*edgenode",
    "conduit_list": {
        "prod": {
            "if_list": [ "10g1", "10g3" ],
            "team_mode": 6
        },
        "public" : {
            "if_list": [ "10g2", "10g4" ],
            "team_mode" : 6
        }
    }
},
{
    "pattern": "team/*/*.*namenode",
    "conduit_list": {
        "prod": {
            "if_list": [ "10g1", "10g3" ],
            "team_mode": 6
        }
    }
},
{
    "pattern": "team/*/*.*filernode",
    "conduit_list": {
        "prod": {
            "if_list": [ "10g1", "10g3" ],
            "team_mode": 6
        }
    }
},
{
    "pattern": "team/*/*.*journalingnode",
    "conduit_list": {
        "prod": {
            "if_list": [ "10g1", "10g3" ],
            "team_mode": 6
        }
    }
},
{
    "pattern": "team/*/*.*datanode",
    "conduit_list": {
        "prod": {
            "if_list": [ "10g1", "10g2" ],
            "team_mode": 6
        }
    }
},
{
    "pattern": ".*/*.*",
    "conduit_list": {
        "prod": {
            "if_list": [ "10g2" ]

```



```

    },
    "admin": {
      "if_list": [ "10g2" ]
    },
    "mgmt": {
      "if_list": [ "1g1" ]
    },
    "public": {
      "if_list": [ "10g3" ]
    }
  }
},
{
  "pattern": "mode/1g_adpt_count/role",
  "conduit_list": {
    "prod": {
      "if_list": [ "10g1" ]
    },
    "admin": {
      "if_list": [ "10g1" ]
    },
    "public": {
      "if_list": [ "10g1" ]
    }
  }
}
],
"networks": {
  "bmc": {
    "conduit": "bmc",
    "vlan": 300,
    "use_vlan": false,
    "add_bridge": false,
    "subnet": "172.16.0.0",
    "netmask": "255.255.255.0",
    "broadcast": "172.16.0.255",
    "router": "172.16.0.1",
    "router_pref": 30,
    "ranges": {
      "router": { "start": "172.16.0.1", "end": "172.16.0.10" },
      "host": { "start": "172.16.0.51", "end": "172.16.0.254" }
    }
  },
  "storage": {
    "conduit": "intf1",
    "vlan": 200,
    "use_vlan": false,
    "add_bridge": false,
    "subnet": "192.168.125.0",
    "netmask": "255.255.255.0",
    "broadcast": "192.168.125.255",
    "ranges": {
      "host": { "start": "192.168.125.10", "end": "192.168.125.239" }
    }
  }
},
"bmc_vlan": {
  "conduit": "mgmt",
  "vlan": 300,
  "use_vlan": false,
  "add_bridge": false,
  "subnet": "172.16.0.0",
  "netmask": "255.255.255.0",
  "broadcast": "172.16.0.255",
  "router": "172.16.0.1",
  "router_pref": 30,
  "ranges": {

```

```

        "host": { "start": "172.16.0.21", "end": "172.16.0.50" }
    },
    "public": {
        "conduit": "public",
        "vlan": 500,
        "use_vlan": false,
        "add_bridge": false,
        "subnet": "192.168.1.0",
        "netmask": "255.255.255.0",
        "broadcast": "192.168.1.255",
        "router": "192.168.1.1",
        "router_pref": 10,
        "ranges": {
            "host": { "start": "192.168.1.10", "end": "192.168.1.25" }
        }
    },
    "admin": {
        "conduit": "prod",
        "vlan": 100,
        "use_vlan": false,
        "add_bridge": false,
        "subnet": "172.16.2.0",
        "netmask": "255.255.254.0",
        "broadcast": "172.16.3.255",
        "router_pref": 20,
        "router": "172.16.2.1",
        "ranges": {
            "host": { "start": "172.16.2.21", "end": "172.16.2.254" },
            "dhcp": { "start": "172.16.3.1", "end": "172.16.3.240" },
            "admin" : { "start": "172.16.2.18", "end": "172.16.2.20" }
        }
    }
},
"deployment": {
    "network": {
        "crowbar-revision": 0,
        "elements": {},
        "element_states": {
            "network": [ "readying", "ready", "applying" ]
        },
        "element_order": [
            [ "network" ]
        ],
        "config": {
            "environment": "network-base-config",
            "mode": "full",
            "transitions": true,
            "transition_list": [ "discovered", "reset", "delete" ]
        }
    }
}
}

```

Appendix C: References

Ghemawat, S. Gobioff, H. and Leung, S.-T. The Google File System. Proceedings of the 19th ACM Symposium on Operating Systems Principles. pp. 29--43. Bolton Landing, NY, USA. 2003. © 2003, ACM.

Borthakur, Dhruba. The Hadoop Distributed File System: Architecture and Design. © 2007, The Apache Software Foundation.

Hadoop DFS User Guide. © 2007, The Apache Software Foundation.

HDFS: Permissions User and Administrator Guide. © 2007, The Apache Software Foundation.

HDFS API Javadoc © 2008, The Apache Software Foundation.

HDFS source code

Pig – <http://developer.yahoo.com/hadoop/tutorial/pigtutorial.html>

Pig – <http://pig.apache.org/docs/r0.6.0/setup.html>

Zookeeper – <http://zookeeper.apache.org/doc/r3.2.2/zookeeperOver.html>

Zookeeper – <https://ccp.cloudera.com/display/CDHDOC/ZooKeeper+Installation>

Zookeeper – http://archive.cloudera.com/cdh/3/zookeeper/zookeeperAdmin.html#sc_zkMultServerSetup

Nagios – <http://www.nagios.org/>

Ganglia – <http://ganglia.sourceforge.net/>

Additional information can be obtained at www.dell.com/hadoop or by e-mailing hadoop@dell.com.

To Learn More

For more information on the Dell | Cloudera Solution, visit:

www.dell.com/hadoop

©2011-2014 Dell Inc. All rights reserved. Trademarks and trade names may be used in this document to refer to either the entities claiming the marks and names or their products. Specifications are correct at date of publication but are subject to availability or change without notice at any time. Dell and its affiliates cannot be responsible for errors or omissions in typography or photography. Dell's Terms and Conditions of Sales and Service apply and are available on request. Dell service offerings do not affect consumer's statutory rights.

Dell, the DELL logo, and the DELL badge, PowerConnect, and PowerVault are trademarks of Dell Inc.